



MOTIVATION

REAL STATE IS AN ECONOMIC DRIVER

Due to its importance in the economy, real state price is an important trend to evaluate the macroeconomics

FINDING NEW COMPONENTS TO CAPTURE REAL STATE VALUE IS KEY TO IMPROVE ACCURACY

As the market gets more complex, it is important to find new components to help accessing house pricing. This dataset brings some non-traditional drivers of value such as neighborhood pollution and amount of industries nearby

A TAGIBLE AND EASY TO WORK DATA

To try out my new Spark capabilities I decided to use a data easy to judge and with a subject I have worked before

CORRELATIONS AND PREDICTION

This data set presents opportunity to get some insights from the correlation among variables and to practice some forecasting

DATA SOURCE

https://www.kaggle.com/c/boston-housingsting





CONTEXT CREATION

```
# Creation of context
# Initialize the spark environment

conf = pyspark.SparkConf().setAppName('odl').setMaster('local')
sc = pyspark.SparkContext(conf=conf)
sqlc = pyspark.sql.SQLContext(sc)
sc
```

PREPARATION OF DATA

```
role = get_execution_role()
bucket='odl-spark19spds6003-001'
data_key = 'mc9bx/mc9bx.csv'
data_location = 's3://{}/{}'.format(bucket, data_key) #s3:// is the way to call a pd.read_csv(data_location) # to read a bucket you need

...

df = sqlc.createDataFrame(pd.read_csv(data_location))
df

DataFrame[ID: bigint, crim: double, zn: double, indus: double, chas: bigint, no x: double, rm: double, age: double, dis: double, rad: bigint, tax: bigint, ptra tio: double, black: double, lstat: double, medv: double]

#Writing to Parquet
parquetPath = '/home/ec2-user/SageMaker/mc9bx/mc9bx-hw3-pqt' # changed the name df.write.parquet(parquetPath)
```

VECTORIZATION

```
# make a user defined function (udf)
sqlc.registerFunction("oneElementVec", lambda d: Vectors.dense([d]), returnType=VectorUDT())
# vectorize the data frames
trainingDF = trainingDF.selectExpr("medv", "oneElementVec(rm) as rm") |
testDF = testDF.selectExpr("medv", "oneElementVec(rm) as rm")
print(testDF.orderBy(testDF.medv.desc()).limit(5))
DataFrame[medv: double, rm: vector]
```

REGRESSION COEFICCIENTS

```
from pyspark.ml.regression import LinearRegression, LinearRegressionModel
lr = LinearRegression()
lrModel = lr.fit(trainingDF)
type(lrModel)
print("Coefficient: " + str(lrModel.coefficients))
print("Intercept: " + str(lrModel.intercept))

Coefficient: [8.730576706815834]
Intercept: -31.6555330439087
```





```
numeric_features = [t[0] for t in df.dtypes if t[1] == 'int' or t[1] == 'double']
sampled_data = df.select(numeric_features).sample(False, 0.8).toPandas()
axs = pd.scatter_matrix(sampled_data, figsize=(10, 10))
n = len(sampled_data.columns)
for i in range(n):
    v = axs[i, 0]
    v.yaxis.label.set_rotation(0)
    v.yaxis.label.set_ha('right')
    v.set_yticks(())
    h = axs[n-1, i]
    h.xaxis.label.set_rotation(90)
    h.set xticks(())
```



